

# Automatic Estimation of FM Synthesis Parameters by Convolutional Neural Network

Sergio Rocha da Silva  
Instituto de Ciência e Tecnologia  
Universidade Federal de São Paulo  
São José dos Campos, São Paulo, Brasil  
sergio.silva16@unifesp.br

Márcio P. Basgalupp  
Instituto de Ciência e Tecnologia  
Universidade Federal de São Paulo  
São José dos Campos, SP, Brasil  
basgalupp@unifesp.br

## Abstract

FM synthesis is one of the most famous techniques for sound synthesis and, although it is an old method, it is still widely used nowadays. It is simple, computationally inexpensive, and capable of producing good results. However, successful use of this type of audio synthesis presents a major challenge: the correct choice of parameters for the synthesizer. This process is not precise, and there is no fully deterministic way to follow. Musicians are usually guided by their intuition and an empirical trial-and-error process. To revive this technique and take advantage of its benefits, the use of CNN models appears as a very promising option, since they can help musicians achieve convincing results, imitating real instruments in a short time and at low cost. Therefore, this work explores this approach by training a CNN model to predict the appropriate parameters that enable an FM synthesizer to resynthesize original audio samples (such as acoustic instruments).

## CCS Concepts

• **Do Not Use This Code → Generate the Correct Terms for Your Paper;** *Generate the Correct Terms for Your Paper;* Generate the Correct Terms for Your Paper; Generate the Correct Terms for Your Paper.

## Keywords

frequency modulation, FM Synthesis, CNN, re-synthesis, musical computation, sound, nsynth, convolutional neural network, NSynth

## ACM Reference Format:

Sergio Rocha da Silva and Márcio P. Basgalupp. 2026. Automatic Estimation of FM Synthesis Parameters by Convolutional Neural Network. In *Proceedings of The 41st ACM/SIGAPP Symposium on Applied Computing (SAC'26)*. ACM, New York, NY, USA, 12 pages. <https://doi.org/XXXXXXX.XXXXXXX>

## 1 Introduction

The main goal of this work is to guide a synthesizer based on frequency modulation (FM) to reproduce arbitrary sounds. A secondary goal is to achieve a process that fully explores the capabilities of the FM synthesizer itself.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

SAC'26, Thessaloniki, Greece

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.  
ACM ISBN 979-X-XXXX-XXXX-X/26/03  
<https://doi.org/XXXXXXX.XXXXXXX>

In other words, simulating the input sound is very important, but the resulting model should also be able to predict parameters for any sound that can be generated with the available resources of the synthesizer. This aligns with one of the main historical advantages of synthesizers: the ability to create sounds that cannot be reproduced by acoustic instruments.

Therefore, in addition to recreating acoustic instrument sounds, the model should be able to approximate any sound sample, assisting future users in adapting the synthesizer to reproduce imaginative timbres with fewer manual parameter adjustments. This perspective also guides the choice of the training dataset, as discussed in the Proposed Approach section.

## 2 Theoretical foundations

### 2.1 Musical Computation

Musical computation is a broad and interdisciplinary field that encompasses all aspects of musical processing through a computer. It is beyond the scope of this article to provide a comprehensive overview of this field; however, for the purposes of this study, it is important to cover the basic elements of sound and the fundamentals of audio digitization.

**2.1.1 Sound elements.** Audio is a mechanical compression wave that travels through the air. This type of wave can be converted into an electrical signal through microphones, which consist of an arrangement of coils and magnets. As a wave, its main elements are:

- **Pitch:** The fundamental frequency of a sound, which determines whether it is low or high. The higher the basis frequency, the higher the pitch.
- **Duration:** The length of time a sound is sustained and can be heard.
- **Loudness (or amplitude):** The intensity or energy level of a sound. In terms of a wave, it corresponds to the amplitude of the sound wave.

It is important to highlight that although sound can be described in terms of the aspects mentioned above, hearing is mainly a cognitive process. Therefore, sound comparison cannot be fully addressed by mathematical tools, leaving a significant space for perceptual comparison and classification (not to mention aspects related to personal preferences).

**2.1.2 Sound timbre.** Another important aspect of acoustic sounds is timbre, which can be understood as the “color of sound.”


From a mathematical point of view, timbre varies according to the harmonics present in a sound sample, or, in other words, according to the set of oscillations that compose the sound.

In the case of acoustic instruments, what happens is that the air vibrations produced by the instrument are not simple pure sinusoids. In fact, when a string vibrates, or when air resonates inside a tube, many parts of the instrument's body vibrate simultaneously, and the frequency with which each part vibrates differs according to its composition, position, and other factors.


The result is that each instrument has a unique sound, especially in the case of handcrafted wooden instruments such as violins or cellos.

To obtain a clear visualization of this aspect, it is common to convert a sound sample from the time domain to the frequency domain using a Fourier Transform technique.

For example, consider the two images below, which compare the same violin audio sample in both the time and frequency domains:



**Figure 1: One second time domain violin sample.**




**Figure 2: The same one second frequency domain violin sample.**

As the figures show, there is a main frequency that defines the pitch of the sound, along with many other frequencies that allow

the ear to identify the instrument and distinguish between different kinds of sounds.

**2.1.3 Audio digitization.** A microphone can convert a pressure wave into an electrical waveform. However, in computational music, the audio signal must be discretized to be represented as a digital audio file. The primary process for digital audio representation is the sampling.

Sampling consists of recording a sequence of wave amplitudes at fixed intervals. For example, if a sound is recorded using a sampling rate of 22,100 Hz, it means that the computer will capture 22,100 audio samples per second.



**Figure 3: Signal sampling. Source: Wikimedia Commons, Sampling (signal), CC BY-SA 4.0. [6]**

This approach is very efficient because human hearing can detect frequencies between 20 Hz and 20,000 Hz. However, for most musical sounds, the maximum frequency is typically below 10 kHz (in fact, even a violin — one of the highest-pitched instruments — reaches only around 3,500 Hz). And according to the Nyquist-Shannon Theorem, to properly sample a periodic signal, the sampling rate must be at least twice the frequency of the highest component of the signal.

Thus, 22,100 Hz can be a good sampling rate for recording music while saving disk space. However, if the goal is to cover the full range of human hearing, higher rates, such as 44,200 Hz, can also be used.

Note that with this process, all sound elements highlighted before, can be recorded and reproduced later.

**2.1.4 Sound synthesis.** After the development of sound recording techniques, musical computation evolved toward the creation of artificial waveforms. Of course, before digital audio synthesis, there were several techniques that allowed analog synthesis. However, the underlying idea is essentially quite similar.

In summary, audio synthesis techniques can be divided into three basic types:

- **Additive synthesis:** Perhaps the most straightforward idea from a mathematical point of view. It consists of generating each frequency component of a desired timbre separately and then summing them. In fact, it is more like a weighted


sum, since each component is weighted according to its contribution to the timbre. A normalization step is commonly applied after the process to avoid amplitude overload.

- **Subtractive synthesis:** The opposite of the previous technique, but historically more common due to its effectiveness and lower cost. In short, it consists of taking a complex waveform, rich in harmonics, and applying various filters to obtain the desired sound. It is more empirical than additive synthesis (which is based on spectral analysis), but simpler to implement and less computationally intensive. Typically, the original waveforms are complex ones such as the sawtooth wave, square wave, or triangular wave.
- **FM synthesis:** Another empirical approach to sound synthesis, which will be addressed in detail in a later section. In short, it consists of altering a simple waveform by applying a high-frequency modulation, not to the generated waveform itself, but directly to the frequency parameter of the carrier oscillator. This kind of distortion of the base frequency produces mirrored harmonics that enrich the timbre of the sound. It is up to the musician to choose the appropriate parameters to produce the desired timbre.

There are many other techniques that could be mentioned; however, these three represent the most direct forms of synthesis, or the most purely mathematical ones, in the sense that they do not require manipulation of pre-recorded audio.

**2.1.5 ADSR Audio envelop.** Another important technique for achieving good results is the so-called “Attack, Decay, Sustain, and Release (ADSR) envelope”, which consists of applying a multiplicative function to the waveform, altering only its amplitude.

The basic idea is to control the sound intensity by simulating the natural behavior of acoustic sounds. Typically, an ADSR envelope is defined as a function with a domain between 0 and 1, which is applied to the generated waveform:



**Figure 4: ADSR linear envelop.**

The ADSR envelope can be understood as an amplitude function:

$$y(t) = e(t) \cdot s(t)$$

Where:

- $s(t)$ : the raw generated sound wave.
- $e(t)$ : the ADSR envelope function (normalized amplitude ranging from 0 to 1).

- $y(t)$ : the resulting sound.

Although real audio samples can show a wide variation of envelope behaviors, it is common practice to use a linear function to achieve the goal, since the human ear is more sensitive to abrupt changes and to relative differences. However, for long time variations, the ear can still perceive the change, which is why some synthesizers allow the use of exponential or logarithmic functions as well.

## 2.2 FM Synthesis

As shown in the previous sections, one of the most popular sound synthesis techniques is Frequency Modulation Synthesis (FM Synthesis), which was introduced by John M. Chowning in his famous paper “The Synthesis of Complex Audio Spectra by Means of Frequency Modulation”, published at the Stanford Artificial Intelligence Laboratory.

Chowning explored the effects of frequency modulation when applied with high frequencies, close to the audible range, instead of the traditional use of inaudible frequencies.

In fact, the concept of frequency modulation was already well known before this paper; however, it had been used mainly in radio transmission of music or to apply slight distortions through low-frequency oscillators (LFOs), producing effects such as vibrato.

As Chowning himself pointed out in his paper, when the technique is applied to vary the frequency of the carrier wave through a high-frequency modulator wave, it “results in a surprising control of audio spectra” (Chowning [1]).

In summary, considering only two oscillators: the main one, called the carrier, and the modulator, so the instantaneous frequency of the output audio can be expressed by the following function:

$$f(t) = f_c + I \cdot \sin(2\pi f_m t)$$

Where:

- $f(t)$ : the instantaneous frequency of the generated sound wave.
- $f_c$ : the original frequency of the carrier wave.
- $I$ : the modulation index, which defines the intensity of the modulation (the greater the value, the more harmonics are generated).
- $f_m$ : the frequency of the modulator wave (which determines the spacing of the harmonics).
- $t$ : the time variable.

The most important point, however, is that this simple manipulation results in symmetric sidebands around the carrier frequency, spaced at multiples of  $f_m$ .

For a more intuitive understanding, if a carrier frequency  $f_m$  is modulated by a frequency  $f_m$ , the resulting spectrum will contain:

$$f_c \pm n f_m$$

For example, by choosing  $f_c = 440$  and  $f_m = 220$ , the resulting wave will contain frequency components such as 220 Hz, 440 Hz, 660 Hz, 880 Hz, and so on, with decreasing amplitudes (the amplitudes follow the Bessel coefficients, but this discussion is beyond the scope of this article).

The last important characteristic of Frequency Modulation Synthesis concerns the ratio between the carrier frequency and the modulator frequency  $f_c/f_m$ , since this ratio directly influences the perceived audio result.

This ratio determines whether the generated harmonics align with the natural harmonic series (widely used by musicians) or not:

- If the ratio is an integer number, the harmonic sidebands (in the spectral view) will appear as integer multiples of  $f_c$ , generating harmonic sounds similar to those of real instruments (consistent with traditional music theory).
- If the ratio is not an integer value, the harmonic sidebands will not align perfectly with  $f_c$ , and the generated sound may appear metallic or dissonant (but, this effect is useful for creating sounds such as bells or for producing special audio effects).

## 2.3 Audio comparison

The task of audio comparison is not straightforward, especially when the goal is to achieve equality in terms of human perception. It can be complex, computationally demanding, and inherently subjective.

However, as described in the section on timbre, one intuitive and efficient way to approach this problem is by comparing the audio in its spectral representation, since timbre is precisely defined by this combination of frequencies.

That said, this approach alone is not sufficient, because the Fourier Transform considers the entire audio signal at once, mixing different parts of a sound sample. For example, if a sample contains three different pitches in sequence, its frequency-domain representation will show them simultaneously, which may give the impression that the timbre is a mixture of all these pitches (similar to a chord rather than a sequence).

In fact, according to Claesson [2], FFT-based spectral comparison is valid but only a part of the solution. He highlights the following metrics:

- **FFT Distance:** This metric consists of computing the Fast Fourier Transform (FFT) of the two samples being compared and then calculating the Euclidean distance between them. The result can be normalized so that the maximum distance is 1 and the minimum is 0.
- **Short Time Fourier Transform (STFT) Distance:** Equivalent to the FFT Distance, but calculated over short, overlapping slices of the sample, which provides time-localized spectral information.
- **Log-Mel Spectrogram Distance:** Similar to the previous metrics, but based on the Log-Mel Spectrogram, which is essentially an adaptation of the FFT spectrum that incorporates human auditory perception. In summary, the Log-Mel Spectrogram is obtained by computing the STFT, mapping the frequencies onto the Mel scale (explained later in this document), and finally projecting the amplitudes onto a logarithmic scale.

In addition to these three metrics, Claesson [2] also recommends the use of Euclidean Distance of the Envelope to compare audio envelopes. This approach can be very useful, but since this work is focused on the similarity between original and synthesized audio

in terms of timbre, the envelope comparison, while important for real-world applications, falls outside the scope of the present study.


**2.3.1 Mel-Scale Representation.** The Mel scale was proposed by Stevens, Volkman, and Newman in 1937 (Claesson [2]), and its main purpose is to reflect the human perception of frequency.

The key difference from the linear scale is that, as the values on the scale increase, the perceptual difference between two adjacent points becomes smaller.

In fact, the Mel scale maps linear frequencies onto a logarithmic scale. A frequency  $f$  is converted to the Mel scale through the following formula:

$$\text{Mel}(f) = 2595 \log_{10} \left( 1 + \frac{f}{700} \right)$$

For a better understanding of this conversion, consider the following graph, which compares the linear hertz scale with the Mel scale:



**Figure 5: Comparison between linear hertz scale and Mel scale (Claesson [2]).**

It is a very useful metric for audio comparison, as it takes into account the human auditory perception of frequency.

## 2.4 Convolutional Neural Networks (CNN)

**2.4.1 2D Convolutional Neural Networks.** In summary, a convolutional neural network is based on the idea of filters.

Before the introduction of CNNs, it was already known that many computer vision problems can be solved through the evaluation of image features, which become more evident depending on the filters applied.

For example, to classify different animals, just the silhouettes can be enough. Of course, the fewer features considered, the higher the chance of mistakes, but, as the popular saying goes: “if it walks like a duck, quacks like a duck, flies like a duck... it’s a duck!”. That’s the basic idea.

So, the CNNs were conceived in such a way that several layers of filters of fixed size (chosen by the network designer) are defined, in addition to some pooling layers (which will be discussed later).

These filter layers are also organized according to the designer’s choices (usually based on many experiments) and can be applied in sequence, in parallel, using different combinations, and so on.


However, it is important to note that the filters serve as the link between the different representations of the input image built within the network.

That is, although the network still takes an image as input, and although this image is processed to extract the various features needed for the problem, there are no dense interconnections between these image representations.

In fact, what interposes between one image representation and another is precisely a CNN filter. And even this filter is not densely connected to the image representations.

In practice, the filter is not connected to any specific part of the image. Instead, it is “slid” across the image, performing a Hadamard multiplication (element-wise product) between the filter and each frame of the image (of the same size), producing, as output, a “new filtered image” (multiplication, element by element, of two matrices of the same size, with a final sum of all elements — and, typically, with the application of an activation function to the result).

Visually, the work can be represented as follows:



**Figure 6: Hadamard product between a filter and an input layer slice (adapted from Rokh et al. [4]).**

And it is precisely from this sliding operation that the network architecture derives its name, since this action corresponds to the mathematical operation of convolution, which can be viewed below:

To an input image  $I$ , a filter (or Kernel)  $K$ , the convolution output  $O$ , in a point  $(i, j)$  is given by:

$$O(i, j) = \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} K(m, n) \times I(i + m, j + n)$$

Where:

- $K(m, n)$ : the filter element on position  $(m, n)$ .
- $I(i + m, j + n)$ : the element of the input on position of current window.
- $M, N$ : the height and width of the filter.
- $O(i, j)$ : the output element on position  $(i, j)$ .

Thus, CNNs can perform image filtering while at the same time reducing the amount of memory required compared to traditional fully connected feed-forward networks (because memory is allocated only for the filters and their results, rather than for each individual interconnection weight).

After extracting the necessary features from the data, a small traditional dense (multi-layer) network is added in sequence to the convolutional (and pooling) layers. This dense network is responsible for identifying the patterns required to predict the target variable, whether for classification or regression tasks.

A CNN is typically divided into two stages. The first stage consists of the application of filters (along with pooling layers), while the second stage applies a fully connected network. In this second stage, the output of the first part undergoes a flattening process so

that it can be used as input to the fully connected layers, which handle the usual tasks of classification or regression.

### 3 Related works

It is not the first work to explore the idea of using a neural network to predict the appropriate parameters for a sound synthesizer. However, the main innovation here is the choice of CNNs to analyze the sound to be imitated. The main reasons are presented below:

- The ability of CNNs to process raw audio inputs.
- The potential of CNNs to capture the temporal structure of sound samples.
- The near real-time performance of CNNs, after training, which can be an advantage for future practical applications.

However, the following works were consulted to support this research and to guide the methodological comparison. The results, in fact, are difficult to compare due to the intrinsically different measurement scales:

Claesson [2] proposed the re-synthesis of instrumental sounds using machine learning models combined with an FM synthesizer. However, Claesson [2] did not use CNNs or the raw signal. Instead, he employed the STFT spectrogram as input (in other words, a frequency-domain representation of the signal) which may simplify the learning process but does not fully exploit the temporal information present in the audio.

Steinmetz et al. [5] introduced DDX7, a differentiable FM synthesizer designed for the re-synthesis of musical instrument sounds. The strength of this work lies in the use of a type of CNN known as a Temporal Convolutional Network (TCN). However, their approach differs from the present work in terms of the chosen target values for the model. DDX7 operates directly on the Mel-spectrum representation, integrating it into the loss function. This was remarkable, as it provided the neural network with a loss function aligned with one of the most common audio comparison metrics. At the same time, this approach requires the synthesizer to be included in the training loop, which is more computationally expensive in terms of both memory and processing. Furthermore, compared with the present work, this research emphasizes the idea of strong generalization by generating a wide variety of parameter combinations through controlled random selection of FM parameters, as explained in the next section.

Finally, Kiranyaz et al. [3] presented a comprehensive survey on 1D convolutional neural networks (CNNs) and their applications, highlighting their effectiveness in sequential data such as audio signals. Although this work does not focus on the same purpose (namely, the re-synthesis of audio samples) it explores various use cases for one-dimensional convolutional neural networks, including audio signal processing. This survey is an important reference that validates the applicability of CNNs to audio signals and their ability to capture the temporal structure of sounds.

### 4 Proposed approach

The overall idea of the project can be summarized as follows:

*Use AI to guide an FM synthesizer in simulating any instrument, or any desired sound.*

Or, in a more detailed explanation:



*Use a Convolutional Neural Network (CNN) to analyze a sound signal sample, and predict the correct parameters for an FM synthesizer to re-synthesize the “same” timbre, enabling the synthesizer to simulate that desired “voice” playing any song.*

However, this idea is still quite abstract and broad. Therefore, consider the following steps to achieve the project’s goal:

- (1) Implement a simple FM synthesizer, which can later be used both to generate the dataset and to test the ability to re-synthesize a sound.
- (2) Generate an audio dataset composed of random samples, each created by randomly selecting FM synthesizer parameters. This essentially records tuples of (parameter values, audio sample) and explores the parameter space on a large scale.
- (3) Train a CNN to predict the correct parameters for the same FM synthesizer, given only an audio sample.
- (4) Evaluate the model’s results using the NSynth dataset to test its generalization capabilities.

Frequency Modulation Synthesis was chosen because it is one of the most common and effective methods for generating audio samples, owing to its simplicity and low computational cost.

This technique became extremely popular in the 1980s, when Yamaha released the Yamaha DX7 synthesizer, and in the 1990s it became a standard for many commercial synthesizers, and it is still used today in the production of electronic music (Claesson [2]).

However, the effective use of this technique depends on advanced knowledge to select the right parameters for generating good results, which can involve thousands of parameters.

Therefore, the ability to simplify parameter selection constitutes a strong justification for this study, and AI models appear as an affordable way to learn the right patterns for controlling an FM synthesizer.

The basic idea is that a model can be produced to guide an FM synthesizer in mimicking any sound, thereby optimizing the effort to compose electronic music.

Furthermore, in terms of computational cost, using an AI model to predict parameters and guiding an FM synthesizer is much more effective than applying a Generative Neural Network to generate an audio. Additionally, a synthesizer can be best applied in real-time composition and electronic musician performance.

## 4.1 Training the CNN

As described in the “Theoretical Background” section, the main characteristic of the sound that should be reproduced to achieve the main work goal is the “timbre”, which is defined as the combination of frequencies present in the sound.

It is clear that techniques such as filtering, Fourier Transform, and etc, could be used to extract the timbre components from an audio signal, but it is not enough to control an FM Synthesizer. In fact, the FM Synthesis controlling is, normally, an experimental process involving trial and error.

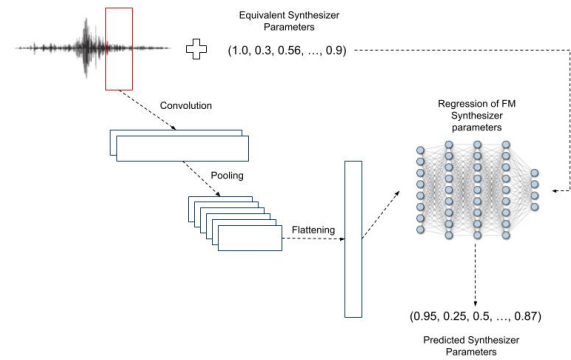
Additionally, as described before, each FM Synthesis operator produces additional frequency components in the signal, and there is no practical method to predict the correct parameter settings, but, as a mathematical function, it is easy to see that this synthesis pattern can be easily approximated using AI models.

Thus, the current work aims to use a CNN model, since it is a very effective method for learning patterns, respecting the internal data structure, and without requiring a lot of preprocessing work. In other words, a CNN can simultaneously learn the time-dependent aspects of an audio signal and can act as a pattern extractor, by applying a sequence of filtering operations, allowing the use of the raw audio format as input.

To train the CNN model, a randomly generated dataset will be used. Its main advantage is that it contains both the audio signal and the corresponding parameters used to generate each sample. In addition, this approach covers a wide range of the parameter space.

Consequently, the training process is straightforward and is expected to yield good performance.

The figure below presents the overall scheme of this training process.



**Figure 7: The overall training scheme.**

**4.1.1 Evaluating Results.** As introduced in the “Theoretical Background” section, the following audio similarity metrics will be used in this work:

- FFT distance: Used to compare the frequency components present in an audio signal.
- STFT distance: Used to compare the variation of frequency components over time.
- Log-mel-spectrogram distance: Used for a similar purpose as FFT distance, but typically models human cognitive pitch perception more accurately.

These metrics will be applied to perform a final test of the model’s generalization capability.

The proposed approach is, after the model training, to use another external dataset, the Google NSynth dataset, which is better described below, and will be used in order to determine whether the model is able to re-synthesize real instruments, following the steps outlined below:

- (1) Train the model with the random generated dataset.
- (2) Use the model to predict the parameters required to re-synthesize the NSynth dataset.
- (3) Re-synthesize the NSynth dataset using the predicted parameters.

- (4) Calculate the audio similarity metrics.
- (5) Evaluate the results.

## 5 Experiments

### 5.1 Setup

*5.1.1 Implemented FM Synthesizer.* There are several well-known FM synthesizers on the market, but most are not open-source. Additionally, commercial synthesizers often have a vast number of parameters that can be used to control the audio production process. For example, Native Instruments' FM8 can use around 1,000 parameters (Claesson [2]), and Teenage Engineering's OP-1 synthesizer can be configured with up to  $10^{76}$  parameters (Claesson [2]).

This incredibly large number of parameters is intended to give users fine-grained control over the generated audio, allowing for a wide variety of pitches (simulating a broad range of instruments and even non-instrument sounds, such as bird songs or car horns, among others).

However, since the scope of this project is not focused on a commercial solution but rather on research purposes, a simple FM synthesizer will suffice.

Therefore, this work will implement a basic FM synthesizer with only 6 operators, inspired by the classic Yamaha DX7, which was one of the most disruptive milestones in the history of electronic music composition. In fact, this simple setup is enough to cover all FM synthesis theories and achieve good performance during AI model evaluation.

In summary, the implemented synthesizer will consist of:

- 5 modulator operators.
- 1 carrier operator.
- 1 ADSR envelope.
- 2 synthesis algorithms (one completely serial, and another with 3 serial operators followed by 2 parallel operators).

The synthesizer will be implemented in Python, using the NumPy library for numerical processing and the SciPy library for audio file generation (in addition to the python-soundfile library).

The figure below presents the FM Synthesizer algorithms:

**[TODO]**

To summarize, the implemented synthesizer renders the audio signal according to the following steps:

- (1) Generate each operator signal by applying it to modulate the phase of the next operator in the sequence. In this implementation, only the phase is varied (similar to the Yamaha DX7), while a full FM synthesis would also vary the frequency. This simplification helps to avoid complex noise-handling issues. All signals are rendered at a 48 kHz sample rate.
- (2) Combine the output of each operator according to the synthesis algorithm (when there are parallel operators, they are summed).
- (3) Apply the ADSR envelope to the signal.
- (4) Downsample the signal from 48 kHz, to 16 kHz (applying also a low-pass filter to avoid aliasing).

And it is important to highlight that the high sample rate of 48 kHz was chosen to allow the FM synthesis to work well enough, since at high frequencies there is a risk of obtaining aliasing and noise (due to the extremely high frequency components generated

by the normal spacing of the FM operators). Furthermore, the low-pass filtering, used to downsample, is also useful to prevent the aliasing problem (after the downsampling process).

*5.1.2 Generating the dataset.* This project will use two datasets, each for a specific purpose: the first is the generated dataset, used for training, and the second is the NSynth dataset, used to evaluate the model's generalization capabilities.

One of the most challenging steps in the KDD process is obtaining and processing data. However, for the purposes of this work, an effective approach can be used. Similar to Claesson's approach (Claesson [2]), since the goal is to achieve good control over the synthesizer parameters, it is sufficient to generate many samples of synthesizer parameters and their respective audio outputs.

Therefore, the main dataset will consist of 5,000 samples generated by the implemented FM synthesizer, with the following characteristics:

- 5,000 samples.
- 4 seconds in duration.
- Monophonic samples (1-channel audio).
- 16 kHz audio sample rate (the synthesizer will generate samples at 48 kHz, but they will be downsampled to 16 kHz, to be compatible with the NSynth dataset).
- Pitches varying between 20 Hz and 6,000 Hz.
- Random ADSR envelopes.

The expectation is to achieve a level of learning where the model can accurately predict the parameters for a wide range of sounds that can be generated by the implemented FM synthesizer.

However, as the main purpose is to re-synthesize real instruments samples, the randomness of this dataset should be controlled to achieve reasonable music property samples, or, in other words, the dataset was constructed respecting the well-known limits and characteristics of real instruments:

- The minimum frequency for the carrier is 20 Hz, and the maximum is 6 kHz (for example, the maximum frequency for a violin is approximately 3 kHz; another important aspect for this maximum choice is that we need to stay under 8 kHz due to the output sample rate, because of the Nyquist frequency). The frequency is chosen by a uniform random distribution.
- The minimum ratio for each modulator is 1/8, and the maximum is 8 (to avoid aliasing in the final signal). Two methods are used to select the ratio: a) with 90% probability, the ratio will be a random uniform choice between discretized values (for example: 1/8, 1/6, ..., 2, 3, 54, ..., 8; the idea is to simulate a real instrument); and b) with 10% probability, the ratio will be a uniform log distribution between 1/8 and 8 (allowing random metallic sounds, like bells, but respecting human auditory perception, where the difference between low frequencies is more perceived than between high frequencies).
- The minimum beta is 0, and the maximum is 8 (also to avoid aliasing in the final signal). The beta is chosen by a uniform random distribution, controlled such that 20% of the values are between 0 and 0.5, 60% are between 0.5 and 3, and 20% are between 3 and 8 (prioritizing rich timbres and avoiding pure sinusoidal sounds, as well as excessive metallic sounds).

**Table 1: Random Uniform ADSR Time Intervals**

Parameter	Interval
Attack	[0.005, 0.5]
Decay	[0.02, 0.6]
Sustain	[0.1, 0.9]
Release	[0.05, 0.8]

- The amplitude is fixed at 1.0, to avoid that different combinations of beta and amplitude generate the same frequency spectrum and the same timbre (prioritizing the CNN learning).

About the ADSR envelop, the following adjustments were made: With these additional adjustments, the generated dataset will be more representative of simulating real instruments.

**5.1.3 External dataset.** In addition to the generated dataset, an external dataset called NSynth, from Google's Magenta project, will be used.

This dataset has a similar configuration to the generated one, and its main purpose is to evaluate the generalization capabilities of the model. This is because it contains many real instrument samples, including both acoustic and electronic instruments (as well as some synthetic sounds generated by professional synthesizers), which the model has never seen before.

Only the test set of the NSynth dataset will be used, with the following characteristics:

- 4,096 samples.
- 4 seconds in duration.
- Monophonic samples (1-channel audio).
- 16 kHz audio sample rate.
- 11 instrument families.
- Pitches ranging from 27.5 Hz to 4,186.01 Hz.
- The envelope includes a sustain of up to 3 seconds and a decay of 1 second.

## 5.2 Training Protocol

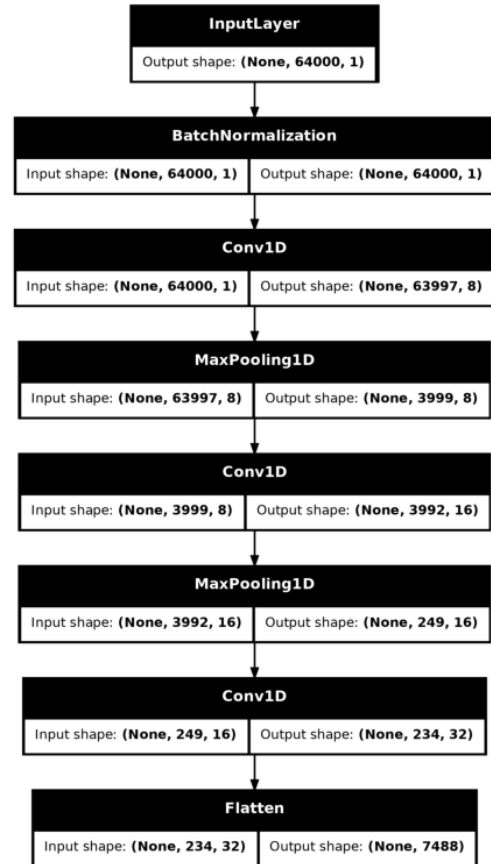
**5.2.1 Environment Settings.** To train the model, the following environment settings were used:

- Python Version: 3.11.2
- Tensorflow Version: 2.19.0
- Keras Version: 3.11.1
- Nvidia-cuda Version: 12.5.82
- Nvidia-cudnn Version: 3.0.75
- Nvidia Driver Version: 535.247.01
- CUDA Version: 12.2
- GPU: NVIDIA GeForce GTX 1650 (4GB memory)

**5.2.2 Train/Test Split.** As each sample in the generated dataset is a complete 4-second sound (containing 64,000 audio samples, which corresponds to 16 kHz multiplied by 4 seconds), a simple 75/25 train-test split was used, resulting in 3,750 audio files for training and 1,250 for testing.

But, from these 3,750 audio files, 20% of them will be used for validation, resulting in a total of 3,000 samples for training and 750 samples for testing.

**5.2.3 CNN Architecture and Hyperparameters.** The model architecture used is as following images, where the first cover the features extraction (CNN), and the second cover the regression (a common MLP):

**Figure 8: The features extraction part (CNN).**

The complete model is simple the sum of the two parts. But it is important to highlight the hyperparameters of each part:

Other important highlights:

- The target variable is normalized to have a mean of 0 and a standard deviation of 1, prior to training.
- The Nadam optimizer was used.
- The Mean Squared Error (MSE) was used as the loss function.
- The recorded metrics were the Mean Absolute Error (MAE), and the Mean Squared Error (MSE).
- The model was trained for 200 epochs, with early stopping enabled using 20 epochs of patience and best weights restoring.
- The model was trained with a batch size of 32.



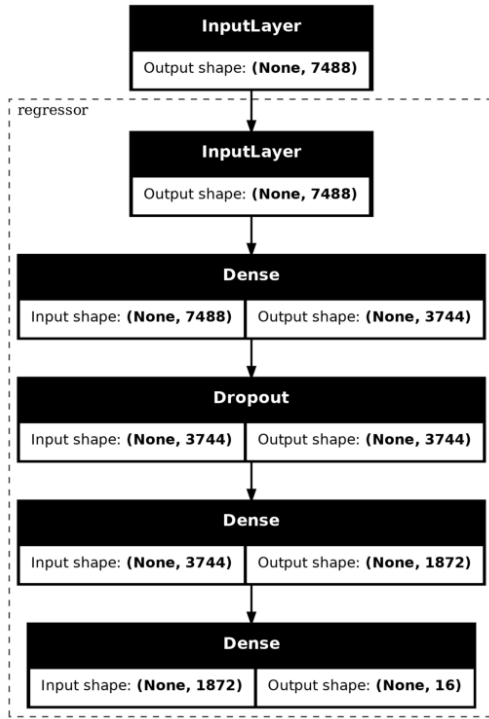


Figure 9: The regressor part (common MLP).

**5.2.4 Hyperparameters Optimization.** After an initial phase of experiments aimed at selecting a generic architecture to achieve the proposed goals, a grid search process was applied to optimize the model's hyperparameters.

This grid search did not differ substantially from the traditional process, but it is important to highlight the search space and the chosen strategy to obtain a good representation of the parameter space without requiring exhaustive experimentation.

The following hyperparameters and their respective options were used:

However, calculating the total number of possible parameter combinations results in the remarkable value of 5,184 options. An exhaustive search over this space would take nearly nine days, given that 186 experiments required approximately seven hours.

To make the process viable, a random shuffle of the parameter search space was applied, producing a new order of experiments that did not require exhausting all options of a given parameter before moving on to the next.

Using this shuffled hyperparameter space, 186 experiments were conducted, representing nearly 3.6% of the total and requiring about seven hours of training.

This approach enabled the representation of a generic sample of the entire search space within a feasible time frame. Since any parameter option could appear in the sampled configurations, the random grid search allowed for faster and broader coverage of the parameter space.

Furthermore, each combination was used to train for 30 epochs with an early stopping patience of 5. All tested hyperparameters

Table 2: Hyperparameters of Features Extraction (CNN part)

Component	Hyperparameter	Value
Input Layer	Input shape	$(input\_len, input\_dims)$
BatchNormalization	–	Applied to input layer
	Filters	8
Conv1D (1)	Kernel size	4
	Strides	1
	Activation	relu
	Regularization	None
MaxPooling1D (1)	Use bias	True
	Pool size	16
	Filters	16
Conv1D (2)	Kernel size	8
	Strides	1
	Activation	relu
	Regularization	None
MaxPooling1D (2)	Use bias	True
	Pool size	16
	Filters	32
Conv1D (3)	Kernel size	16
	Strides	1
	Activation	relu
	Regularization	None
MaxPooling1D (3)	Use bias	True
	Pool size	16
Flatten	–	After Conv1D layers
BatchNormalization	–	Applied after Flatten

Table 3: Hyperparameters of the Regressor Network

Component	Hyperparameter	Value
Input Layer	Input shape	$(input\_dims)$
	Units	$input\_dims/2$
Dense (1)	Activation	relu
	Use bias	True
	Kernel regularizer	l2
Dropout	Rate	0.5
	Units	$input\_dims/4$
Dense (2)	Activation	relu
	Use bias	True
	Units	$output\_dims$
Dense (Output)	Activation	None
	Use bias	True
Optimizer	–	Nadam
Loss	–	MSE
Metrics	–	MAE, MSE

and their respective error metric values were recorded using the MLflow tool, which enabled more accurate comparisons among the top solutions. This, in turn, helped guide the selection of good hyperparameters and supported a more targeted exhaustive training phase, already informed by the most successful configurations.


**Table 4: Hyperparameters and Available Options**

Hyperparameter	Options
Activation	gelu, swish, elu, relu, selu, silu
Bias (CNN)	True, False
Kernel Regularizer (CNN)	None, l1, l2
Bias (Regressor)	True, False
Dropout (Regressor)	0, 0.2, 0.3, 0.5
Kernel Regularizer (Regressor)	None, l1, l2
Optimizer	AdamW, RMSprop, Adam,
–	Nadam, Adagrad, Adamax

**Table 5: Training and Validation Metrics by Epoch**

Epoch	Loss	Train		Validation		
		MAE	MSE	Loss	MAE	MSE
70	0.603	0.527	0.545	0.622	0.517	0.562
71	0.601	0.526	0.542	0.637	0.520	0.569
72	0.605	0.526	0.545	0.621	0.516	0.565
73	0.602	0.524	0.541	0.625	0.522	0.568
74	0.611	0.528	0.548	0.625	0.520	0.563

To provide a clearer understanding of this point, consider the following experiment comparison graph:

**Figure 10: MLflow grid search results showing the best error metrics (based on the test set).**

11

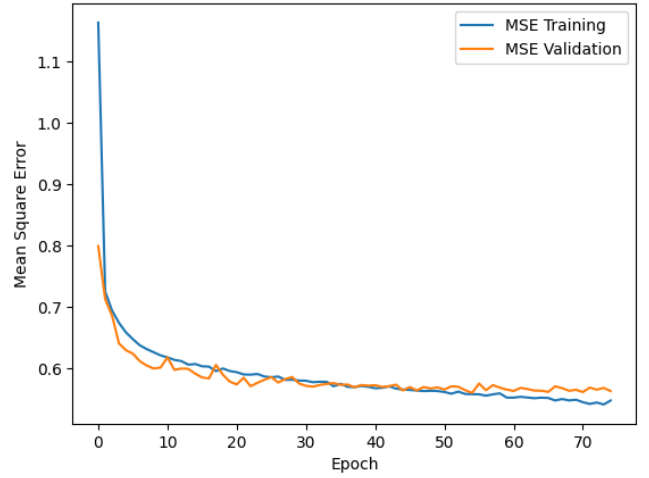
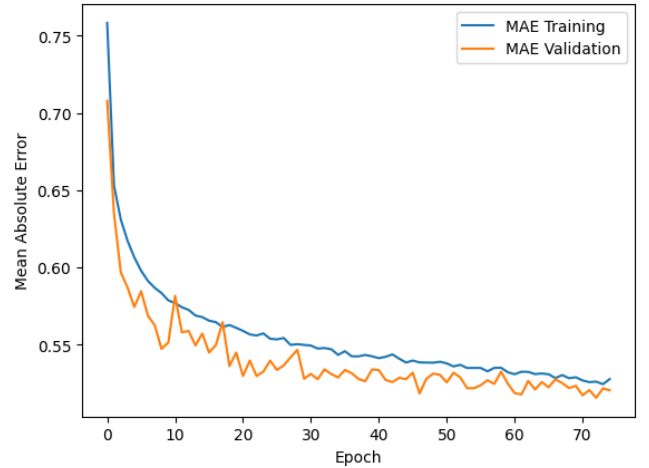
Note in the figure above that the following configurations appear successful, based on the three error metrics:

- **Activation function:** RELU
- **CNN Bias:** True
- **Regressor Bias:** True
- **Regressor Dropout:** 0.5
- **CNN Kernel Regularizer:** None
- **Regressor Kernel Regularizer:** L2

It should be noted that the two best configurations, considering the MAE metric, were tested, and the second one (with respect to MAE) showed the best results without signs of overfitting.

## 5.3 Results

**5.3.1 Training/Validation Results.** The training and validation results are shown in the following table, which summarizes the last five epochs of training:

**Figure 11: Evolution of MSE during training.****Figure 12: Evolution of MAE during training.****Table 6: Evaluation Metrics on Test Set**

Metric	Value
RMSE	37.551
MSE	1410.097
MAE	6.488

**5.3.2 Test Results.** To compute the test results, the predicted values, from the splitted test portions, was transformed back to the original space, and the following metrics were calculated:

**5.3.3 Evaluation Results.** As explained in the section about the proposed approach, after the CNN training, to evaluate the model's ability to re-synthesize real instruments, the NSynth dataset was used, in order to calculate specific metrics focused on the capacity of the FM synthesizer to recreate the same timbres.

**Table 7: Comparison of Means (Resynthesis vs Baseline, with Improvement %)**

Metric	Resynthesis	Baseline	Improvement (%)
FFT	12649.920	33508.780	62.27
STFT	2366.424	7290.362	67.54
Log Mel (raw)	138.732	184.872	24.96
Log Mel (normalized)	0.00860	0.01146	24.93

**Table 8: Comparison of Standard Deviations (Resynthesis vs Baseline)**

Metric	Resynthesis	Baseline
FFT	6208.149	1887.364
STFT	1132.152	434.867
Log Mel (raw)	47.013	30.136
Log Mel (normalized)	0.00292	0.00187

To achieve this, the model was applied to the NSynth test dataset, and the predicted parameters were used with the implemented FM synthesizer. Then, each signal output was compared with the respective original signal, calculating the proposed metrics:

- FFT distance
- STFT distance
- Log-mel-spectrogram distance

However, before using the predicted parameters as input to the synthesizer, a simple function was applied to approximate each ratio parameter to the nearest discrete ratio, according to the common behavior of musical harmonics (the same set of ratios used to randomize 90% of the generated dataset). Furthermore, some clipping operations were also applied to enforce the same limits imposed on the randomly generated dataset, such as maximum and minimum values for beta parameters.

In order to achieve a good evaluation, the proposed metrics were also used to compare each NSynth audio with a baseline, which consists of a pure sinusoidal signal, with a frequency of 440 Hz (i.e., the A4 note, which is used as a reference to tune the majority of instruments).


Then, the results can be seen in the following:

Note that the Log Mel metric was compared, including a normalized version, which is obtained by dividing the raw Log Mel value by the product of the number of Mel bands and the time window quantity.

The normalized Log Mel is proposed because it is better suited for comparing results with other works. For example, authors such as [TODO] achieved good results with values between 0.005 and 0.02, which indicates that the current work result is within a satisfactory range.

Another important insight can be obtained from a histogram representation of the Log-Mel distances:

Furthermore, to facilitate a better understanding, the metrics were also compared by grouping the dataset by instrument family. This allows for the evaluation of the model's performance on different classes of sounds:

**Figure 13: Distribution of Normalized Log-Mel Distance.****Table 9: Comparison of Means by Instrument Group**

Group	FFT	STFT	Log Mel (raw)	Log Mel (norm)
bass	11232.735	2144.318	135.610	0.00841
brass	14371.068	2682.950	167.617	0.01039
flute	14611.274	2872.685	141.819	0.00879
guitar	9692.000	1867.466	119.853	0.00743
keyboard	9619.258	1772.917	119.529	0.00741
mallet	8085.799	1632.454	120.610	0.00748
organ	18360.101	3453.125	166.433	0.01032
reed	12740.034	2267.079	140.701	0.00872
string	12883.688	2390.651	139.365	0.00864
vocal	31028.909	5271.256	212.646	0.01319

**Table 10: Comparison of Standard Deviations by Instrument Group**

Group	FFT	STFT	Log Mel (raw)	Log Mel (norm)
bass	5271.640	931.050	43.094	0.00267
brass	3934.388	949.146	50.102	0.00311
flute	3711.774	769.482	36.122	0.00224
guitar	3805.303	716.923	38.010	0.00236
keyboard	3024.448	507.672	34.145	0.00212
mallet	2304.329	437.872	38.022	0.00236
organ	3418.068	818.277	47.884	0.00297
reed	4205.592	905.286	46.032	0.00285
string	4696.608	889.743	35.951	0.00223
vocal	8218.038	1672.940	56.507	0.00350

And as can be seen, on average, the top three models performed best for the keyboard, guitar and mallet families (in this order).

## 6 Conclusion

As proposed in the introduction of this work, there are two main goals to achieve with this technique: the first is to re-synthesize the sounds of real instruments, and the second is to explore the parameter space in a randomized way, testing whether the model

can generalize the results. This generalization is a strong indication that the model can guide the synthesizer to reproduce arbitrary sounds.

As shown in the results, the CNN model demonstrated good generalization ability (as observed in the test results over the randomly generated dataset) and also showed strong performance in the re-synthesis of real instruments (as evidenced by the audio comparison metrics between the NSynth dataset and the corresponding re-synthesized audio).

Furthermore, the results indicate that the model is not extremely biased toward specific instrument families. As presented in the previous section, the normalized log-mel metric remained within a reasonable range across all instrument families, and the difference between the best and worst cases was not excessively high.

The proposed approach clearly shows that both pillars of the initial hypothesis were valid:

- CNNs are highly effective for feature extraction from instrument audio samples.
- The approach of generating a randomized dataset, aiming to cover a wide area of the parameter space, was efficient in achieving good control over a specific FM synthesizer.

However, this work also has clear limitations that should be highlighted, which at the same time represent opportunities for future research:

- **Generic dataset:** Adopt a new dataset composed of generic audio samples, not limited to musical instruments (e.g., bells, horns, bird calls, etc.), which could serve as inspiration for musicians.
- **Broader ADSR envelope variations:** Since the main goal was related to real instrument re-synthesis, and the test dataset (NSynth) is composed of such sounds, only a common linear ADSR envelope was implemented. While sufficient for this purpose, it would certainly not be adequate for a more generic dataset.
- **Exploring more FM synthesis algorithms:** During the experiments, only two FM algorithms were evaluated, and only one produced good results (the one with three serial operators and two parallel operators, which was used for all reported results). However, professional FM synthesizers typically allow arbitrary operator configurations, which could be explored further.

## 6.1 Future Work

In order to highlight future research opportunities, the following directions appear most promising:

- Explore the use of more generic datasets.
- Implement and evaluate a wider range of ADSR envelopes and FM algorithms.
- Investigate less straightforward neural network architectures, testing whether feature extraction through core representation (e.g., autoencoder networks) could be effective for generic datasets.
- Combine approaches from related works: for example, using the STFT of each sample as input to the regression stage of the network (Claesson [2]) alongside CNN-based feature extraction from raw signals, or applying the log-mel metric

directly in the loss function, as in DDX7 (Steinmetz et al. [5]).

- Experiment with AutoML techniques after CNN feature extraction, as a way to improve the regression stage of the model. In this context, the autoencoder hypothesis could also be used to pre-train the CNN independently.
- Adopt a professional FM synthesizer implementation (with the necessary permissions) to re-train the entire model.

## Acknowledgments

This work is supported by *Coordenação de Aperfeiçoamento de Pessoal de Nível Superior* (CAPES), *Conselho Nacional de Desenvolvimento Científico e Tecnológico* (CNPq), grants 406417/2022-9, 102475/2024-5, and 312209/2022-3, and *Fundação de Amparo à Pesquisa do Estado de São Paulo* (FAPESP), grant 2020/09835-1 (CPA IARA).

## References

- [1] John M. Chowning. 1973. The synthesis of complex audio spectra by means of frequency modulation. *Journal of the Audio Engineering Society*, 21, 7, 526–534.
- [2] Philip Claesson. 2019. *Re-synthesis of instrumental sounds with Machine Learning and a Frequency Modulation synthesizer*. Master's thesis. KTH Royal Institute of Technology, School of Electrical Engineering and Computer Science.
- [3] Serkan Kiranyaz, Onur Avci, Osama Abdeljaber, Turker Ince, Moncef Gabbouj, and Daniel J. Inman. 2021. 1d convolutional neural networks and applications—a survey. *Journal of Computer Science and Technology*, 36, 1, 1–25.
- [4] Babak Rokh, Ali Azarpeyvand, and Alireza Khanteymoori. 2022. Hadamard product of a filter and a part of layer's input. Accessed: 2025-05-28. ResearchGate. [https://www.researchgate.net/figure/Hadamard-product-of-a-filter-and-a-part-of-layers-input\\_fig16\\_360668087](https://www.researchgate.net/figure/Hadamard-product-of-a-filter-and-a-part-of-layers-input_fig16_360668087).
- [5] Christian J. Steinmetz, Jordi Pons, Dmitry Bogdanov, and Xavier Serra. 2022. Ddx7: differentiable fm synthesis of musical instrument sounds. In *Proceedings of the 23rd International Society for Music Information Retrieval Conference (ISMIR)*.
- [6] Wikipedia contributors. 2024. Sampling (signal). Accessed: 2025-10-04. Wikimedia Foundation. [https://pt.wikipedia.org/wiki/Amostragem\\_de\\_sinal](https://pt.wikipedia.org/wiki/Amostragem_de_sinal).